

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, VotingClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
from imblearn.over_sampling import SMOTE
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: # Load the dataset
df = pd.read_csv('C:/Users/srirk/Downloads/archive (1)/WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
In [3]: # Display basic information about the dataset
print("Basic Information about the Dataset:")
print(df.info())
print("\nFirst 5 Rows of the Dataset:")
print(df.head())
```

Basic Information about the Dataset:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 7043 entries, 0 to 7042

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64
6	PhoneService	7043 non-null	object
7	MultipleLines	7043 non-null	object
8	InternetService	7043 non-null	object
9	OnlineSecurity	7043 non-null	object
10	OnlineBackup	7043 non-null	object
11	DeviceProtection	7043 non-null	object
12	TechSupport	7043 non-null	object
13	StreamingTV	7043 non-null	object
14	StreamingMovies	7043 non-null	object
15	Contract	7043 non-null	object
16	PaperlessBilling	7043 non-null	object
17	PaymentMethod	7043 non-null	object
18	MonthlyCharges	7043 non-null	float64
19	TotalCharges	7043 non-null	object
20	Churn	7043 non-null	object

dtypes: float64(1), int64(2), object(18)

memory usage: 1.1+ MB

None

First 5 Rows of the Dataset:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	
4	9237-HQITU	Female	0	No	No	2	Yes	

	MultipleLines	InternetService	OnlineSecurity	... DeviceProtection	\
0	No phone service	DSL	No	...	No
1	No	DSL	Yes	...	Yes
2	No	DSL	Yes	...	No
3	No phone service	DSL	Yes	...	Yes
4	No	Fiber optic	No	...	No

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	
2	No	No	No	Month-to-month	Yes	
3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	

	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

In [4]:

```
# Handling missing values
df.replace(' ', np.nan, inplace=True)
print("\nMissing Values before Imputation:")
print(df.isnull().sum())
imputer = SimpleImputer(strategy='median')
df['TotalCharges'] = imputer.fit_transform(df[['TotalCharges']])

print("\nMissing Values after Imputation:")
print(df.isnull().sum())
```

Missing Values before Imputation:

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64
```

Missing Values after Imputation:

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
```

```

Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64

```

In [5]:

```

# Encode categorical variables
categorical_features = df.select_dtypes(include=['object']).columns
for col in categorical_features:
    if col != 'customerID':
        df[col] = LabelEncoder().fit_transform(df[col])

df.drop('customerID', axis=1, inplace=True)
print("\nDataset after Encoding Categorical Variables and Dropping customerID Column")
print(df.head())

```

Dataset after Encoding Categorical Variables and Dropping customerID Column:

```

gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0      0             0      1            0        1            0
1      1             0      0            0        34            1
2      1             0      0            0         2            1
3      1             0      0            0        45            0
4      0             0      0            0         2            1

MultipleLines  InternetService  OnlineSecurity  OnlineBackup  \
0             1                0                0            2
1             0                0                2            0
2             0                0                2            2
3             1                0                2            0
4             0                1                0            0

DeviceProtection  TechSupport  StreamingTV  StreamingMovies  Contract  \
0                0            0            0            0            0
1                2            0            0            0            1
2                0            0            0            0            0
3                2            2            0            0            1
4                0            0            0            0            0

PaperlessBilling  PaymentMethod  MonthlyCharges  TotalCharges  Churn
0                1              2           29.85       29.85      0
1                0              3           56.95      1889.50      0
2                1              3           53.85       108.15      1
3                0              0           42.30      1840.75      0
4                1              2           70.70       151.65      1


```

In [6]:

```


# Visualize the distribution of churn
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='Churn')
plt.title('Distribution of Churn')
plt.show()

```




In [7]:


```
# Correlation heatmap
plt.figure(figsize=(14, 10))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```



```
In [8]: # Visualize the distribution of features
plt.figure(figsize=(8, 6))
sns.histplot(df['tenure'], bins=20, kde=True)
plt.title('Distribution of Tenure')
plt.show()
```



```
In [9]: plt.figure(figsize=(8, 6))
sns.histplot(df['MonthlyCharges'], bins=20, kde=True)
plt.title('Distribution of Monthly Charges')
plt.show()
```



```
In [10]: # Feature scaling
scaler = StandardScaler()
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns
df[numerical_features] = scaler.fit_transform(df[numerical_features])

In [11]: # Splitting data into features and target
X = df.drop('Churn', axis=1)
y = df['Churn']

In [12]: # Handle class imbalance with SMOTE
smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)

In [13]: # Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2, ran

In [14]: # Initialize models
log_reg = LogisticRegression(random_state=42)
rf_clf = RandomForestClassifier(random_state=42)
gb_clf = GradientBoostingClassifier(random_state=42)

In [15]: # Hyperparameter tuning using Grid Search
param_grid_log_reg = {
    'C': [0.01, 0.1, 1, 10, 100],
    'solver': ['lbfgs', 'liblinear']
}

param_grid_rf = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

param_grid_gb = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 4, 5]
}

grid_search_log_reg = GridSearchCV(LogisticRegression(random_state=42), param_grid_log_reg)
grid_search_rf = GridSearchCV(RandomForestClassifier(random_state=42), param_grid_rf)
grid_search_gb = GridSearchCV(GradientBoostingClassifier(random_state=42), param_grid_gb)

In [16]: # Train models with best parameters
grid_search_log_reg.fit(X_train, y_train)
grid_search_rf.fit(X_train, y_train)
grid_search_gb.fit(X_train, y_train)

best_log_reg = grid_search_log_reg.best_estimator_
best_rf_clf = grid_search_rf.best_estimator_
best_gb_clf = grid_search_gb.best_estimator_

# Predict on the test set
y_pred_log_reg = best_log_reg.predict(X_test)
y_pred_rf_clf = best_rf_clf.predict(X_test)
```

```
y_pred_gb_clf = best_gb_clf.predict(X_test)

# Function to print evaluation metrics
def evaluate_model(y_true, y_pred, model_name):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    conf_matrix = confusion_matrix(y_true, y_pred)

    print(f"Evaluation Metrics for {model_name}:")
    print(f"Accuracy: {accuracy:.2f}")
    print(f"Precision: {precision:.2f}")
    print(f"Recall: {recall:.2f}")
    print(f"F1 Score: {f1:.2f}")
    print(f"Confusion Matrix:\n{conf_matrix}\n")

    return accuracy
```

```
In [17]: # Evaluate Logistic Regression
acc_log_reg = evaluate_model(y_test, y_pred_log_reg, "Logistic Regression")
```

```
Evaluation Metrics for Logistic Regression:
Accuracy: 0.80
Precision: 0.77
Recall: 0.86
F1 Score: 0.81
Confusion Matrix:
[[749 272]
 [150 899]]
```

```
In [18]: # Evaluate Random Forest Classifier
acc_rf_clf = evaluate_model(y_test, y_pred_rf_clf, "Random Forest Classifier")
```

```
Evaluation Metrics for Random Forest Classifier:
Accuracy: 0.85
Precision: 0.83
Recall: 0.88
F1 Score: 0.86
Confusion Matrix:
[[838 183]
 [124 925]]
```

```
In [19]: # Evaluate Gradient Boosting Classifier
acc_gb_clf = evaluate_model(y_test, y_pred_gb_clf, "Gradient Boosting Classifier")
```

```
Evaluation Metrics for Gradient Boosting Classifier:
Accuracy: 0.84
Precision: 0.84
Recall: 0.86
F1 Score: 0.85
Confusion Matrix:
[[847 174]
 [152 897]]
```

```
In [20]: # Ensemble method using Voting Classifier
voting_clf = VotingClassifier(estimators=[
    ('log_reg', best_log_reg),
```

```
    ('rf_clf', best_rf_clf),  
    ('gb_clf', best_gb_clf)  
], voting='soft')  
  
voting_clf.fit(X_train, y_train)  
y_pred_voting = voting_clf.predict(X_test)
```

In [21]:

```
# Evaluate Voting Classifier  
acc_voting = evaluate_model(y_test, y_pred_voting, "Voting Classifier")
```

Evaluation Metrics for Voting Classifier:

Accuracy: 0.85

Precision: 0.82

Recall: 0.90

F1 Score: 0.86

Confusion Matrix:

```
[[818 203]
```

```
 [109 940]]
```

In [22]:

```
# Print accuracies for each model  
print("Model Accuracies:")  
print(f"Logistic Regression Accuracy: {acc_log_reg:.2f}")  
print(f"Random Forest Classifier Accuracy: {acc_rf_clf:.2f}")  
print(f"Gradient Boosting Classifier Accuracy: {acc_gb_clf:.2f}")  
print(f"Voting Classifier Accuracy: {acc_voting:.2f}")
```

Model Accuracies:

Logistic Regression Accuracy: 0.80

Random Forest Classifier Accuracy: 0.85

Gradient Boosting Classifier Accuracy: 0.84

Voting Classifier Accuracy: 0.85